



## Code Vectorization & Tools

Bertrand Putigny

July 6, 2017



# Outline

## Introduction

Vectorization: the Hardware  
Code Vectorization

## Case Study

## Conclusion

# Vectorization (with OpenMP4)

- ◇ What is vectorization?
- ◇ Do I need vectorization
- ◇ How to "vectorize" code?

# Vector Instruction

## SIMD: Single Instruction Multiple Data

A0	A1	A2	A3
+			
B0	B1	B2	B3
=			
A0+B0	A1+B1	A2+B2	A3+B3

# SIMD Instructions Sets

- ◇ SSE: 128bits
  - 2 double precision reals
  - 4 single precision reals
- ◇ AVX: 256bits
  - 4 double precision reals
  - 8 single precision reals
- ◇ coming up: AVX-512: 512bits

**SIMD is here to stay:**

Trends:

- ◇ larger vectors
- ◇ more instructions (FMA, gather...)

⇒ need to optimize code for SIMD

# Vectorization: Using vector instructions

- ◇ automatic code vectorization (compiler)
  - ◇ hand written code
    - intrinsics
    - assembly
- ⇒ poor portability, hard to write, hard to read

# Vectorization: Using vector instructions

- ◇ automatic code vectorization (compiler)
  - ◇ hand written code
    - intrinsics
    - assembly
- ⇒ poor portability, hard to write, hard to read

## Solution:

Understanding basics of compiler vectorization:

- ◇ code transformation
- ◇ why vectorization can fail
- ◇ help compiler in such cases

# Automatic Code Vectorization

## Code transformation:

Do the same thing "differently":

- ◇ keep the same semantic
- ◇ different code versions
- ◇ can be done at several level
  - source code level (source to source compilers)
  - intermediate representation (most of the time)
  - instruction level

## Code transformation examples:

- ◇ instruction scheduling (optimize ILP, at assembly level)
- ◇ scalar promotion (IR level)

```
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        A[i][j] = (1/(double) i) * A[i][j];  
    }  
}
```

- ◇ loop tiling (cache access optimization, most of the time by hand)



# Automatic Code Vectorization

## Code Transformation:

1. rely on loop unrolling
2. turn set of instructions (scalar) into a single vector instruction

## Original code:

```
for(i=0; i<SIZE; i++) {  
    y[i] = x[i] + y[i];  
}
```

## 1. Unrolled loop:

```
// peeling (if need be)  
for(i=0; i<SIZE-SIZE%4; i+=4) {  
    y[i] = x[i] + y[i];  
    y[i+1] = x[i+1] + y[i+1];  
    y[i+2] = x[i+2] + y[i+2];  
    y[i+3] = x[i+3] + y[i+3];  
}  
// remainder...
```

## 2. Vectorized pseudo-code:

```
for(i=0; i<SIZE-SIZE%4; i+=4) {  
    y[i:i+3] = x[i:i+3] + y[i:i+3];  
}  
// remainder...
```

# Factor Affecting Code Vectorization: Trip Count

## Scalar code:

```
for(i=0; i<7; i++) {  
    y[i] = x[i] + y[i];  
}
```

$\approx 7$  cycles

## Vectorized:

```
for(i=0; i<4; i+=4) {  
    y[i:i+3] = x[i:i+3] + y[i:i+3];  
}  
y[4] = x[4] + y[4];  
y[5] = x[5] + y[5];  
y[6] = x[6] + y[6];
```

$\approx 4$  cycles

## Vectorized with padding:

```
for(i=0; i<8; i+=4) {  
    y[i:i+3] = x[i:i+3] + y[i:i+3];  
}
```

$\approx 2$  cycles

# Factor Affecting Code Vectorization: Dependencies


## Loop-carried data dependencies:

- ◇ cannot be vectorized:

```
for(i=1; i<SIZE; i++) {  
    y[i] = y[i-1] - y[i];  
}
```

- ◇ can be vectorized if vector length  $\leq 4$ :

```
for(i=4; i<SIZE; i++) {  
    y[i] = y[i-4] - y[i];  
}
```



⇒ use OpenMP 4.0 `pragma omp simd safelen(n)`

# Factor Affecting Code Vectorization: Aliasing

## Pointer Aliasing:

```
void foo(double *x, double *y, int n) {  
    for(i=0; i<n; i++) {  
        x[i] = y[i] - x[i];  
    }  
}
```

```
void bar() {  
    foo(x, x+1, n-1);  
}
```


⇒ use compiler `-fno-alias` option (if you do not use aliasing)

⇒ `#pragma omp simd` (code level, scope)

# Factor Affecting Code Vectorization: Data Layout


## Poor memory access:

```
struct coord {  
    double x;  
    double y;  
};  
  
for(i=0; i<n; i++) {  
    points[i].x += v.x;  
    points[i].y += v.y;  
}
```



## Optimal memory access:

```
struct coord {  
    double *x;  
    double *y;  
};  
  
for(i=0; i<n; i++) {  
    points.x[i] += v.x[0];  
    points.y[i] += v.y[0];  
}
```



# Factor Affecting Code Vectorization: Control Flow

## Conditionals:

```
for(i=0; i<n; i++) {  
    if (x[i] > threshold) {  
        x[i] = y[i];  
    }  
}
```

## Function calls:

```
for(i=0; i<n; i++) {  
    x[i] = f(y[i]);  
}
```

⇒ use OpenMP 4.0 `pragma omp declare simd`

# Factor Affecting Code Vectorization: Reduction

## Sum:

```
r = .0;
for(i=0; i<n; i++) {
    r += x[i];
}
```

⇒ use pragma omp reduction(+: r)

## Code Samples:

`https://github.com/bputigny/T8.A02`



# Conclusion

## Code optimization

multiprocessor  $\times$  mono processor optimization

## Vectorization

- ◇ rely on compiler vectorization
- ◇ help compiler with:
  - OpenMP 4
  - memory layout
  - code transformation (if need be)

## Tools

Help yourself with:

- ◇ VTune: multiprocessor profiling
- ◇ Intel Advisor: vectorization advisor